



Module 2

Introduction


We have seen how to **build and optimize deep feedforward architectures** consisting of linear & non-linear (e.g. ReLU) layers

- ⬢ This can be generalized to **arbitrary computation graphs**
- ⬢ **Backpropagation and automatic differentiation** can be used to optimize all parameters via **gradient descent**



But layers don't have to be fully connected – **other connectivity structures are possible**

For images, it makes sense for output nodes to **consider only small patches of inputs**




This operation, which can be another layer in the neural network, can be formalized as **convolution operations**

$$K = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$




This new convolution layer can take **any input 3D tensor (say, RGB image)** and **output another similarly-shaped output**

- Where **high values represent strong responses** of the learned features



Convolution layers can be combined with non-linear and **pooling layers** which reduce the dimensionality of the data



Adding Other Layers




Image



Conv +
Non-Linear
Layer



Pooling
Layer



Conv +
Non-Linear
Layer



Fully
Connected
Layers

Alternating Convolution and Pooling

These architectures have existed **since 1980s**





Image Credit: Yann LeCun, Kevin Murphy

LeNet Architecture

But have become **deeper and more complex**



FC

Conv
1x1+1(S)

MaxPool
3x3+1(S)


SoftmaxActivation

From: Szegedy et al. Going deeper with convolutions


ResNet Architecture



Layer 1




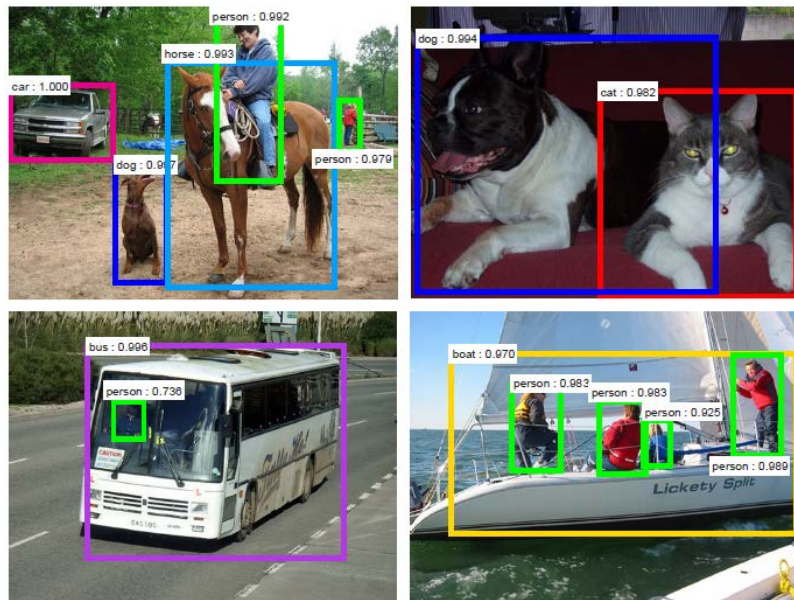
Layer 2



From: Zeiler & Fergus, Visualizing and Understanding Convolutional Networks

Visualization of Learned Features

More complex tasks, like **object detection**, show the real power of deep learning



From: Ren et al., Faster R-CNN

Object Detection Architectures